

```
1 import keras
2 keras.__version__

1 from keras import layers
2 from keras import models
3
4 model = models.Sequential()
5 model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
6 model.add(layers.MaxPooling2D((2, 2)))
7 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
8 model.add(layers.MaxPooling2D((2, 2)))
9 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

Let's display the architecture of our convnet so far:

```
1 model.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
conv2d_15 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_10 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_16 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_11 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_17 (Conv2D)	(None, 3, 3, 64)	36928

=====
Total params: 55,744
Trainable params: 55,744
Non-trainable params: 0
=====

You can see above that the output of every `Conv2D` and `MaxPooling2D` layer is a 3D tensor of shape (height, width, channels) . The width and height dimensions tend to shrink as we go deeper in the network. The number of channels is controlled by the first argument passed to the `Conv2D` layers (e.g. 32 or 64).

The next step would be to feed our last output tensor (of shape (3, 3, 64)) into a densely-connected classifier network like those you are already familiar with: a stack of `Dense` layers. These classifiers process vectors, which are 1D, whereas our current output is a 3D tensor. So first, we will have to flatten our 3D outputs to 1D, and then add a few `Dense` layers on top:

```
1 model.add(layers.Flatten())
2 model.add(layers.Dense(64, activation='relu'))
3 model.add(layers.Dense(10, activation='softmax'))
```

We are going to do 10-way classification, so we use a final layer with 10 outputs and a softmax activation. Now here's what our network looks like:

```
1 model.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
conv2d_15 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_10 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_16 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_11 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_17 (Conv2D)	(None, 3, 3, 64)	36928
flatten_7 (Flatten)	(None, 576)	0
dense_14 (Dense)	(None, 64)	36928

dense_15 (Dense) (None, 10) 650

```
=====
Total params: 93,322
Trainable params: 93,322
Non-trainable params: 0
=====
```

As you can see, our (3, 3, 64) outputs were flattened into vectors of shape (576,), before going through two Dense layers.

Now, let's train our convnet on the MNIST digits. We will reuse a lot of the code we have already covered in the MNIST example from Chapter 2.

```
1 from keras.datasets import mnist
2 from tensorflow.keras.utils import to_categorical
3
4 (train_images, train_labels), (test_images, test_labels) = mnist.load_data()
5
6 train_images = train_images.reshape((60000, 28, 28, 1))
7 train_images = train_images.astype('float32') / 255
8
9 test_images = test_images.reshape((10000, 28, 28, 1))
10 test_images = test_images.astype('float32') / 255
11
12 train_labels = to_categorical(train_labels)
13 test_labels = to_categorical(test_labels)

1 model.compile(optimizer='adam',
2               loss='categorical_crossentropy',
3               metrics=['accuracy'])
4 model.fit(train_images, train_labels, epochs=7, batch_size=100)

Epoch 1/7
600/600 [=====] - 43s 70ms/step - loss: 0.0939 - accuracy: 0.9714
Epoch 2/7
600/600 [=====] - 42s 70ms/step - loss: 0.0532 - accuracy: 0.9837
Epoch 3/7
600/600 [=====] - 42s 70ms/step - loss: 0.0370 - accuracy: 0.9880
Epoch 4/7
600/600 [=====] - 42s 70ms/step - loss: 0.0285 - accuracy: 0.9909
Epoch 5/7
600/600 [=====] - 42s 70ms/step - loss: 0.0230 - accuracy: 0.9926
Epoch 6/7
600/600 [=====] - 42s 69ms/step - loss: 0.0180 - accuracy: 0.9940
Epoch 7/7
600/600 [=====] - 42s 69ms/step - loss: 0.0163 - accuracy: 0.9942
<keras.callbacks.History at 0x7f548b433a90>
```

Let's evaluate the model on the test data:

```
1 test_loss, test_acc = model.evaluate(test_images, test_labels)

☐ 313/313 [=====] - 3s 8ms/step - loss: 0.0379 - accuracy: 0.9890

1 test_acc

0.9890000224113464

1 test_loss

0.037858348339796066
```

While our densely-connected network from Chapter 2 had a test accuracy of 97.8%, our basic convnet has a test accuracy of 99.3%: we decreased our error rate by 68% (relative). Not bad!

✓ 0s completed at 10:22 PM

